

# Cycle 08 Homework Notes

**Name:** A Chetan Varma

**Date:** 22-07-2025

---

## Auto Scaling Group (ASG)

---

### Part 1: Hands-On ASG Lab

#### 1. Launch an EC2 Instance

- AMI: Ubuntu (latest LTS)
- Instance type: t2.micro
- Security Group: Allow HTTP (port 80) and SSH (port 22)
- Commands to install Apache and serve a test page:

```
bash
CopyEdit
sudo apt update
sudo apt install apache2 -y
echo "<h1>Apache Test Page</h1>" | sudo tee /var/www/html/index.html
sudo systemctl enable apache2
sudo systemctl start apache2
```

- Validate by visiting <http://<EC2-Public-IP>>
- 

#### 2. Create a Custom AMI

- Stop the EC2 instance
- Go to EC2 → Actions → Image → Create Image
- Name: `My_Apache_AMI`

- Wait for image status to become available
- 

### **3. Create a Launch Template**

- Use the newly created AMI
  - Instance type: t2.micro
  - Security Group: Same as before (HTTP + SSH)
  - Keep default settings for other options
  - Save and create the launch template
- 

### **4. Set Up an Auto Scaling Group (ASG)**

- Use the launch template
  - Choose at least 2 subnets in different Availability Zones
  - ASG configuration:
    - Minimum capacity: 1
    - Desired capacity: 1
    - Maximum capacity: 3
  - Health check type: EC2
  - Health check grace period: 60 seconds
- 

### **5. Configure Scaling Policy**

- Add Target Tracking Policy:
    - Target value: 60 percent CPU utilization
    - Cooldown period: 120 seconds
  - Optionally, install CloudWatch Agent for more metrics
- 

### **6. Test Scaling Behavior**

#### **a. Spike CPU Load**

- SSH into the running instance:

```
bash
CopyEdit
sudo apt install stress -y
stress -c 2
```

- Monitor EC2 console → ASG should launch a new instance if CPU > 60% for 2 minutes

## b. Terminate Instance Manually

- Terminate an instance from ASG manually
- ASG should automatically launch a replacement to maintain desired count

---

## Part 2: Advanced ASG Experiment

### 1. Scale Based on HTTP Requests (ALB + ASG)

- Create an Application Load Balancer (ALB)
  - At least 2 subnets
  - Target group: Register ASG
- Attach ALB to the ASG
- Create scaling policy:
  - Metric: ALB RequestCountPerTarget
  - Target value: Example 50 requests per minute
- Send test traffic using tools like Apache Benchmark:

```
bash
CopyEdit
```

```
ab -n 1000 -c 100 http://<ALB-DNS-Name>/
```

## 2. Multi-AMI Testing

- Create a second EC2 instance with Nginx:

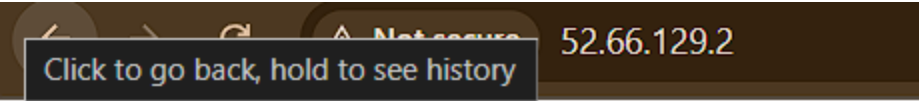
```
bash
CopyEdit
sudo apt update
sudo apt install nginx -y
echo "Nginx Test Page" | sudo tee /var/www/html/index.html
```

- Stop instance and create AMI: `My_Nginx_AMI`
- Create another Launch Template using this new AMI
- Create a separate ASG to compare Apache vs Nginx scaling behavior

### SCREENSHOTS:

```
ubuntu@ip-172-31-9-253:~$ sudo -i
root@ip-172-31-9-253:~# cat > /var/www/html/index.html
WELCOME TO THE PAGE
```

```
ubuntu@ip-172-31-9-253:~$ history
1  sudo apt-get update
2  sudo apt-get install apache2
3  sudo systemctl enable apache2
4  sudo systemctl restart apache2
5  sudo systemctl status apache2
6  sudo systemctl status apache2
```



WELCOME TO THE PAGE

Instances (1/1) Info

Last updated 1 minute ago

Connect

Instance state ▼

Actions ▲

Launch instances

Find Instance by attribute or tag (case-sensitive)

All states ▼

<input checked="" type="checkbox"/>	Name <a href="#">🔗</a>	Instance ID	Instance state ▼	Instance type ▼	Status check	Alarm s
<input checked="" type="checkbox"/>	Image-Autosc...	i-00529d0d57e703277	Running <a href="#">🔍</a>	t2.micro	2/2 checks pass <a href="#">🔍</a>	<a href="#">View al</a>

Create image

Instance diagnostics

Instance settings ▶

Networking ▶

Security ▶

Image and templates ▶

Amazon Machine Images (AMIs) (1) Info

Recycle Bin

EC2 Image Builder

Actions ▼

Launch instance from AMI

Owned by me ▼

Find AMI by attribute or tag

< 1 >

⚙️

<input type="checkbox"/>	Name <a href="#">🔗</a>	AMI name	AMI ID	Source	Owner	Visibility
<input type="checkbox"/>	AMI-Image		ami-01a7df6a215de1f69	845958739988/AMI-Image	845958739988	Private

Instances (1/1) Info

Last updated less than a minute ago

Connect

Instance state ▼

Actions ▲

Launch instances

Find Instance by attribute or tag (case-sensitive)

All states ▼

<input checked="" type="checkbox"/>	Name <a href="#">🔗</a>	Instance ID	Instance state ▼	Instance type ▼	Status check	Alarm s
<input checked="" type="checkbox"/>	Image-Autosc...	i-00529d0d57e703277	Running <a href="#">🔍</a>	t2.micro	2/2 checks pass <a href="#">🔍</a>	<a href="#">View al</a>

Create image

Create template from instance

Instance diagnostics

Instance settings ▶

Networking ▶

Security ▶

Image and templates ▶

Monitor and troubleshoot ▶

Connect

Instance state ▼

Actions ▼

Launch instances

▲

All states ▼

Launch instances

Launch instance from template

Migrate a server

▼	Status check	Alarm status
---	--------------	--------------

### ▼ Application and OS Images (Amazon Machine Image) [Info](#)

An AMI contains the operating system, application server, and applications for your instance. If you don't see a suitable AMI below, use the search field or choose [Browse more AMIs](#).

Q Search our full catalog including 1000s of application and OS images

AMI from catalog   Recents   **My AMIs**   Quick Start

☐ Don't include in launch template

☒ Owned by me

☐ Shared with me

[Browse more AMIs](#)  
Including AMIs from AWS, Marketplace and the Community

#### Amazon Machine Image (AMI)

AMI-Image  
ami-01a7df6a215de1f69  
2025-07-22T09:42:15.000Z   Virtualization: hvm   ENA enabled: true   Root device type: ebs   Boot mode: uefi-preferred

Description

## Instances (2) [Info](#)

Last updated less than a minute ago [Refresh](#) [Control](#)

Q Find Instance by attribute or tag (case-sensitive)



<input type="checkbox"/>	Name <a href="#">✎</a> ▼	Instance ID	Instance state ▼	Instance type ▼
<input type="checkbox"/>	Image2-Autos...	i-0da7bd3534c1ad3d4	Running <a href="#">⚙</a> <a href="#">🔍</a>	t2.micro
<input type="checkbox"/>	Image-Autosc...	i-00529d0d57e703277	Running <a href="#">⚙</a> <a href="#">🔍</a>	t2.micro

← → ↻ [⚠ Not secure](#) 3.108.215.41

WELCOME TO THE PAGE

## 3d4 (Image2-Autoscaling) [Info](#)

### Public IPv4 address

 3.108.215.41 | [open address](#) 

### > Create Auto Scaling group


For most applications, you can use multiple Availability Zones and let EC2 instances span across multiple Availability Zones. The default VPC and default subnets are suitable for getting started.

#### VPC

Choose the VPC that defines the virtual network for your Auto Scaling group.

vpc-00dcef4d69d6b4f5d  
172.31.0.0/16 Default



[Create a VPC](#) 

#### Availability Zones and subnets

Define which Availability Zones and subnets your Auto Scaling group can use in the

Select Availability Zones and subnets



aps1-az1 (ap-south-1a) | subnet-  
00bb156050e0ef090  
172.31.32.0/20 Default




aps1-az2 (ap-south-1c) | subnet-  
090e19b6a34692968  
172.31.16.0/20 Default



aps1-az3 (ap-south-1b) | subnet-  
0a14d154b1b33442a  
172.31.0.0/20 Default



[Create a subnet](#) 

## Create Auto Scaling group

### Health checks

Health checks increase availability by replacing unhealthy instances. If at least one health check fails, instance replacement occurs.

#### EC2 health checks

[i Always enabled](#)

#### Additional health check types - *optional* | [Info](#)

- ☒ Turn on Elastic Load Balancing health checks  
Elastic Load Balancing monitors whether instances are healthy. If an instance is unhealthy, it can replace it on its next periodic check.
- ☒ Turn on VPC Lattice health checks  
VPC Lattice can monitor whether instances are available. If an instance is unhealthy, it replaces it after its next periodic check.
- ☐ Turn on Amazon EBS health checks  
EBS monitors whether an instance's root volume or attached volumes are healthy. If a volume is unhealthy, it replaces the instance on its next periodic health check.

#### Health check grace period | [Info](#)

This time period delays the first health check until your instances enter a non-running state.

seconds



## Create Auto Scaling group

Scaling

### Group size [Info](#)

Set the initial size of the Auto Scaling group. After creating the group, you can change the size manually or by using automatic scaling.

### Desired capacity type

Choose the unit of measurement for the desired capacity value. vCPUs and Memory are configured with a set of instance attributes.

Units (number of instances) ▼

### Desired capacity

Specify your group size.

1

### Scaling [Info](#)

You can resize your Auto Scaling group manually or automatically.

### Scaling limits

Set limits on how much your desired capacity can be increased or decreased.

#### Min desired capacity

1

Equal or less than desired capacity

#### Max desired capacity

3

Equal or greater than desired capacity

## Create Auto Scaling group

### Choose whether to use a target tracking policy [Info](#)

You can set up other metric-based scaling policies and scheduled scaling after creating your Auto Scaling group.

☐ No scaling policies

Your Auto Scaling group will remain at its initial size and will not dynamically resize to meet demand.

☒ Target tracking scaling policy

Choose a CloudWatch metric and target value and let the scaling policy adjust the desired capacity in proportion to the metric's value.

### Scaling policy name

Target Tracking Policy

### Metric type [Info](#)

Monitored metric that determines if resource utilization is too low or high. If using EC2 metrics, consider enabling detailed monitoring for better scaling performance.

Average CPU utilization

### Target value

50

### Instance warmup [Info](#)

60 seconds

☐ Disable scale in to create only a scale-out policy

### Tags (1)

Edit

Key	Value	Tag new instances
AUTOSCALING-IMAGES		Yes

## Auto Scaling groups (1) [Info](#)

less than a minute ago

[Launch configurations](#)

[Launch templates](#)

Search your Auto Scaling groups

<input type="checkbox"/>	Name	Launch template/configuration	Instances	Status	Desired capacity	Min	Max
<input type="checkbox"/>	<a href="#">Autoscaling</a>	<a href="#">Template</a>   Version Default	1	-	1	1	3

## Instances (1/4) [Info](#)

less than a minute ago

[Connect](#)

[Instance state](#)

[Actions](#)

[Launch instances](#)

Find Instance by attribute or tag (case-sensitive)

All states

< 1 >

<input checked="" type="checkbox"/>	Name	Instance ID	Instance state	Instance type	Status check	Alarm status	Availability Zone	Public IP
<input checked="" type="checkbox"/>	Image-Autosc...	i-06e79828dd3cd9ea8	Running	t2.micro	2/2 checks passed	<a href="#">View alarms</a>	ap-south-1a	ec2-13-1...
<input type="checkbox"/>	Image-Autosc...	i-0e8b71cfb3bae70d	Terminated	t2.micro	-	<a href="#">View alarms</a>	ap-south-1a	-

```
root@ip-172-31-42-141:~# top
```

```

top - 10:08:23 up 6 min, 1 user, load average: 0.00, 0.03, 0.01
Tasks: 105 total, 1 running, 104 sleeping, 0 stopped, 0 zombie
%Cpu(s): 0.0 us, 0.0 sy, 0.0 ni, 100.0 id, 0.0 wa, 0.0 hi, 0.0 si, 0.0 st
MiB Mem : 957.4 total, 457.4 free, 339.6 used, 316.6 buff/cache
MiB Swap: 0.0 total, 0.0 free, 0.0 used. 617.8 avail Mem

```

PID	USER	PR	NI	VIRT	RES	SHR	S	%CPU	%MEM	TIME+	COMMAND
1	root	20	0	22144	13300	9460	S	0.0	1.4	0:04.13	systemd
2	root	20	0	0	0	0	S	0.0	0.0	0:00.00	kthreadd
3	root	20	0	0	0	0	S	0.0	0.0	0:00.00	pool_workqueue_release
4	root	0	-20	0	0	0	I	0.0	0.0	0:00.00	kworker/R-rcu_g
5	root	0	-20	0	0	0	I	0.0	0.0	0:00.00	kworker/R-rcu_p
6	root	0	-20	0	0	0	I	0.0	0.0	0:00.00	kworker/R-slub
7	root	0	-20	0	0	0	I	0.0	0.0	0:00.00	kworker/R-netns
8	root	20	0	0	0	0	I	0.0	0.0	0:00.03	kworker/0:0-cgroup_destroy
10	root	0	-20	0	0	0	I	0.0	0.0	0:00.00	kworker/0:0H-events_highpri
11	root	20	0	0	0	0	I	0.0	0.0	0:00.07	kworker/u30:0-events_power_efficient
12	root	0	-20	0	0	0	I	0.0	0.0	0:00.00	kworker/R-mm_pe
13	root	20	0	0	0	0	I	0.0	0.0	0:00.00	rcu_tasks_rude_kthread
14	root	20	0	0	0	0	I	0.0	0.0	0:00.00	rcu_tasks_trace_kthread
15	root	20	0	0	0	0	S	0.0	0.0	0:00.06	ksoftirqd/0
16	root	20	0	0	0	0	I	0.0	0.0	0:00.13	rcu_sched
17	root	rt	0	0	0	0	S	0.0	0.0	0:00.00	migration/0
18	root	-51	0	0	0	0	S	0.0	0.0	0:00.00	idle_inject/0
19	root	20	0	0	0	0	S	0.0	0.0	0:00.00	cpuhp/0
20	root	20	0	0	0	0	S	0.0	0.0	0:00.00	kdevtmpfs
21	root	0	-20	0	0	0	I	0.0	0.0	0:00.00	kworker/R-inet
22	root	20	0	0	0	0	I	0.0	0.0	0:00.03	kworker/u30:1-events_unbound

```

Last login: Tue Jul 22 10:07:02 2025 from 13.233.177.5
root@ip-172-31-42-141:~# dd if=/dev/zero of=/dev/null &
[1] 2433
root@ip-172-31-42-141:~# dd if=/dev/zero of=/dev/null &
[2] 2434
root@ip-172-31-42-141:~# dd if=/dev/zero of=/dev/null &
[3] 2435
root@ip-172-31-42-141:~# dd if=/dev/zero of=/dev/null &
[4] 2436
root@ip-172-31-42-141:~# dd if=/dev/zero of=/dev/null &
[5] 2437
root@ip-172-31-42-141:~# dd if=/dev/zero of=/dev/null &
[6] 2438
root@ip-172-31-42-141:~# dd if=/dev/zero of=/dev/null &
[7] 2439
root@ip-172-31-42-141:~# dd if=/dev/zero of=/dev/null &

```

```

top - 10:22:27 up 20 min, 1 user, load average: 94.40, 55.48, 24.62
Tasks: 204 total, 101 running, 103 sleeping, 0 stopped, 0 zombie
%Cpu(s): 62.4 us, 37.6 sy, 0.0 ni, 0.0 id, 0.0 wa, 0.0 hi, 0.0 si, 0.0 st
MiB Mem : 957.4 total, 459.9 free, 334.0 used, 320.0 buff/cache
MiB Swap: 0.0 total, 0.0 free, 0.0 used. 623.4 avail Mem

```

<input type="checkbox"/>	Name	Instance ID	Instance state	Instance type	Status check	Alarm status	Availability Zone	Public IP
<input type="checkbox"/>	Image-Autosc...	i-0d3447846115f0b6c	Running	t2.micro	2/2 checks passed	<a href="#">View alarms +</a>	ap-south-1b	ec2-65-0-
<input type="checkbox"/>	Image-Autosc...	i-06e79828dd3cd9ea8	Running	t2.micro	2/2 checks passed	<a href="#">View alarms +</a>	ap-south-1a	ec2-13-1-
<input type="checkbox"/>	Image-Autosc...	i-05004bde8a33159	Running	t2.micro	Initializing	<a href="#">View alarms +</a>	ap-south-1a	ec2-3-6-5

	Failed	<p>Launching a new EC2 instance. Status Reason: Your requested instance type (t2.micro) is not supported in your requested Availability Zone (ap-south-1c). Please retry your request by not specifying an Availability Zone or choosing ap-south-1a, ap-south-1b. Launching EC2 instance failed.</p>	<p>At 2025-07-22T10:21:49Z a monitor alarm TargetTracking-Autoscaling-AlarmHigh42c9-b79f-abff31b43a4d in state ALARM triggered policy Target Tracking Policy of capacity from 2 to 3. At 2025-07-22T10:22:02Z an instance was started in response between desired and actual capacity, increasing the capacity from 2 to 3.</p>
	Successful	<p>Launching a new EC2 instance: i-0d3447846115f0b6c</p>	<p>At 2025-07-22T10:19:49Z a monitor alarm TargetTracking-Autoscaling-AlarmHigh42c9-b79f-abff31b43a4d in state ALARM triggered policy Target Tracking Policy of capacity from 1 to 2. At 2025-07-22T10:19:58Z an instance was started in response between desired and actual capacity, increasing the capacity from 1 to 2.</p>
	Successful	<p>Launching a new EC2 instance: i-06e79828dd3cd9ea8</p>	<p>At 2025-07-22T10:01:13Z an instance was launched in response to an unhealthy instance being replaced.</p>
	Successful	<p>Terminating EC2 instance: i-0e8b71fcfb3bae70d</p>	<p>At 2025-07-22T10:01:13Z an instance was taken out of service in response to an EC2 instance indicating it has been terminated or stopped.</p>
	Successful	<p>Launching a new EC2 instance: i-0e8b71fcfb3bae70d</p>	<p>At 2025-07-22T09:59:18Z an instance was started in response to a difference between desired and actual capacity, increasing the capacity from 0 to 1.</p>

Desired capacity	Min	Max	Actual capacity
3	1	3	3

## Part 3: Troubleshooting Tasks

### 1. Simulate Failed Health Check

- SSH into instance and stop Apache:

```
bash
CopyEdit
sudo systemctl stop apache2
```

- Within a minute, ASG should detect health check failure and terminate the instance
  - A new healthy instance should be launched automatically
- 

## 2. Analyze CloudWatch Metrics

- Go to CloudWatch → Metrics → EC2 → Per-Instance Metrics
  - Look for:
    - CPUUtilization
    - GroupDesiredCapacity
    - GroupInServiceInstances
    - GroupTerminatingInstances
  - Verify spikes during stress test and scaling events
- 

## EOD Report Suggestions

- Document screenshots:
  - Apache test page
  - AMI creation
  - Launch template
  - ASG settings
  - Scaling events in EC2 and CloudWatch
- Learning Outcomes:

- Understood autoscaling fundamentals: health checks, templates, and policies
- Practiced scaling based on CPU and HTTP requests
- Learned troubleshooting using service stop and CloudWatch metrics